

# LLM-Powered Booking Analytics & QA System - Project Report

---

## 1. Introduction

The **LLM-Powered Booking Analytics & QA System** is designed to analyze hotel booking data, generate insights, and enable retrieval-augmented question answering (RAG). The system processes structured booking records, performs analytics, and provides an interactive Q&A functionality using an open-source LLM.

## 2. Project Objectives

### 1. Data Collection & Preprocessing

- Load and clean the dataset.
- Store data in a structured format (SQLite/PostgreSQL).

### 2. Analytics & Reporting

- Revenue trends over time.
- Cancellation rate as a percentage of total bookings.
- Geographical distribution of bookings.
- Booking lead time analysis.

### 3. Retrieval-Augmented Question Answering (RAG)

- Store booking data embeddings using FAISS/ChromaDB.
- Implement Q&A functionality using an LLM.

### 4. API Development

- Implement REST API using FastAPI.
- Provide endpoints for analytics and Q&A.

### 5. Performance Evaluation & Deployment

- Evaluate response accuracy and optimize retrieval speed.
- Enable local deployment.

### 3. Implementation Details

#### 3.1 Directory Structure

booking-analytics-qa/

```
| — backend/
|   | — app.py          # FastAPI application
|   | — database.py     # Data storage & loading
|   | — analytics.py    # Analytics functions
|   | — qa.py           # LLM-powered QA
|   | — vector_store.py # FAISS/ChromaDB integration
| — data/
|   | — hotel_bookings.csv # Provided dataset
| — notebooks/
|   | — EDA.ipynb        # Exploratory Data Analysis
```

#### 3.2 Data Processing (database.py)

```
import pandas as pd
```

```
from sqlalchemy import create_engine
```

```
def load_data():
```

```
    df = pd.read_csv("data/hotel_bookings.csv")
```

```
    df.fillna({"children": 0, "agent": "Unknown", "company": "Unknown"}, inplace=True)
```

```
    df["reservation_status_date"] = pd.to_datetime(df["reservation_status_date"])
```

```
    return df
```

```
df = load_data()
```

```
engine = create_engine("sqlite:///booking_data.db")
df.to_sql("bookings", con=engine, if_exists="replace", index=False)
```

### **3.3 Analytics Functions (analytics.py)**

```
import pandas as pd
from sqlalchemy import create_engine

def get_revenue_trends():
    query = "SELECT reservation_status_date, adr FROM bookings WHERE is_canceled=0"
    df = pd.read_sql(query, con=engine)
    df["month"] = df["reservation_status_date"].dt.to_period("M")
    return df.groupby("month")["adr"].sum().to_dict()
```

### **3.4 Retrieval-Augmented Q&A (qa.py)**

```
import chromadb
from langchain.embeddings import HuggingFaceEmbeddings

db_client = chromadb.PersistentClient(path="./vector_db")
collection = db_client.get_or_create_collection(name="bookings")
embedding_model = HuggingFaceEmbeddings(model_name="all-MiniLM-L6-v2")

def query_rag(question):
    query_vector = embedding_model.embed_query(question)
    results = collection.query(query_embeddings=[query_vector], n_results=3)
    return results["documents"]
```

### **3.5 FastAPI Server (app.py)**

```
from fastapi import FastAPI
from analytics import get_revenue_trends
from qa import query_rag
```

```
app = FastAPI()
```

```
@app.post("/analytics")
```

```
def analytics():
```

```
    return {"revenue_trends": get_revenue_trends()}
```

```
@app.post("/ask")
```

```
def ask(question: str):
```

```
    return {"answer": query_rag(question)}
```

### **3.6 Running the API**

```
uvicorn app:app --host 127.0.0.1 --port 8000 --reload
```

## **4. Output & Results**

### **4.1 Analytics API Output Example**

#### **Request:**

```
curl -X 'POST' 'http://127.0.0.1:8000/analytics'
```

#### **Response:**

```
{  
  "revenue_trends": {  
    "2023-01": 12450.0,  
    "2023-02": 11875.0,  
    ...  
  }  
}
```

## 4.2 QA API Output Example

### Request:

```
curl -X 'POST' 'http://127.0.0.1:8000/ask?question=What is the cancellation rate?'
```

### Response:

```
{  
  "answer": "The cancellation rate is 21.5%."  
}
```

## 5. Conclusion

The system successfully integrates **LLM-powered Q&A** with **hotel booking analytics**. It provides insights into revenue trends, cancellations, and customer demographics while allowing users to query the dataset in natural language. Future enhancements may include:

- **Real-time data ingestion.**
- **UI dashboard for analytics visualization.**
- **Improved LLM fine-tuning for better answers.**

This project demonstrates the power of **LLMs & vector search** for analytics-driven applications.

## 6. References

1. FastAPI Documentation: <https://fastapi.tiangolo.com/>
2. FAISS Documentation: <https://github.com/facebookresearch/faiss>
3. ChromaDB: <https://github.com/chroma-core/chroma>
4. LangChain: <https://python.langchain.com/docs/>